# Computer Architecture

## Assignment-1

## Cache Replacement Policies Analysis

The objective of this report is to examine the cache performance of two benchmark programs **test_fmath** and **anagram** - under various cache replacement policies (LRU and MRU) using an out-of-order execution model. This analysis assesses cache hit/miss behavior, execution cycles, and overall efficiency to determine the effects of the replacement strategies.

**Benchmarks**:

- **test_fmath** (floating-point heavy computation)
- **anagram** (string/character manipulation workload)

**Cache Replacement Policies**:

- **LRU (Least Recently Used)**
- **MRU (Most Recently Used)**

**Cumulative Results**

**Alpha ISA**

```
mak@Maahi-VivoBook-S15:~/assgn/simplesim-3.0/scripts$ python3 alpha_eval
.py
Metric               LRU             MRU
------------------------------------------------
sim_IPC              0.4124          0.3535
sim_num_insn         4855            4855
sim_num_refs         1824            1824
sim_elapsed_time     1               1
dl1.miss_rate        0.0797          0.1037
il1.miss_rate        0.0409          0.0559
dl2.miss_rate        0.9902          0.8000
sim_cycle            11772           13735
```

**Pisa ISA**

```
mak@Maahi-VivoBook-S15:~/assgn/simplesim-3.0/scripts$ python3 pisa_eval.
py
Metric               LRU             MRU
------------------------------------------------
sim_IPC              0.9450          0.4412
sim_num_insn         43493           43493
sim_num_refs         12339           12339
sim_elapsed_time     1               1
dl1.miss_rate        0.0204          0.0279
il1.miss_rate        0.0162          0.0598
dl2.miss_rate        0.7716          0.7092
sim_cycle            46023           98588
```

**Key Findings**

**1. Execution Time (Cycles)**

- **Alpha_anagram:**
  - LRU : 11,772 cycles
  - MRU : 13,735 cycles
  - LRU is 16.7% faster than MRU
- **PISA_test-fmath:**
  - LRU : 46,023 cycles
  - MRU : 98,588 cycles
  - LRU is 114.2% faster than MRU

**2. Instructions Per Cycle (IPC)**

- **Alpha_anagram:**
  - LRU : 0.4124
  - MRU : 0.3535
  - LRU shows 16.7% higher IPC
- **PISA_test-fmath:**
  - LRU : 0.9450
  - MRU : 0.4412
  - LRU shows 114.2% higher IPC

**3. Cache Miss Rates**

**L1 Data Cache Miss Rate:**

- **Alpha_anagram**: LRU (7.97%) vs MRU (10.37%) - LRU has 30.1% lower miss rate
- **PISA_test-fmath**: LRU (2.04%) vs MRU (2.79%) - LRU has 36.8% lower miss rate

**L1 Instruction Cache Miss Rate:**

- **Alpha_anagram**: LRU (4.09%) vs MRU (5.59%) - LRU has 36.7% lower miss rate
- **PISA_test-fmath**: LRU (1.62%) vs MRU (5.98%) - LRU has 269.1% lower miss rate

**L2 Data Cache Miss Rate:**

- **Alpha_anagram**: LRU (99.02%) vs MRU (80.00%) - MRU has 23.8% lower miss rate
- **PISA_test-fmath**: LRU (77.16%) vs MRU (70.92%) - MRU has 8.8% lower miss rate

**Summary of Key Findings:**

**Why MRU Fails ?**

MRU removes the most recently used block when the buffer is full but as per the program flow statistics that word is most likely to be re-used again which is why MRU has higher miss rates in L1 cache compared to LRU.

**Why LRU is Good ?**

Unlike, MRU, LRU believes in **temporal locality** which means that when a block of memory is accessed at a certain point in time then that same block more likely to be accessed in the near future this helps LRU stay ahead of MRU. Therefore any critical instructions which are removed by MRU might be causing us many stalls but those won't necessarily present in case of LRU leading us to **better performance.**

## Graphs

- **L1 Data Cache Miss Rate**

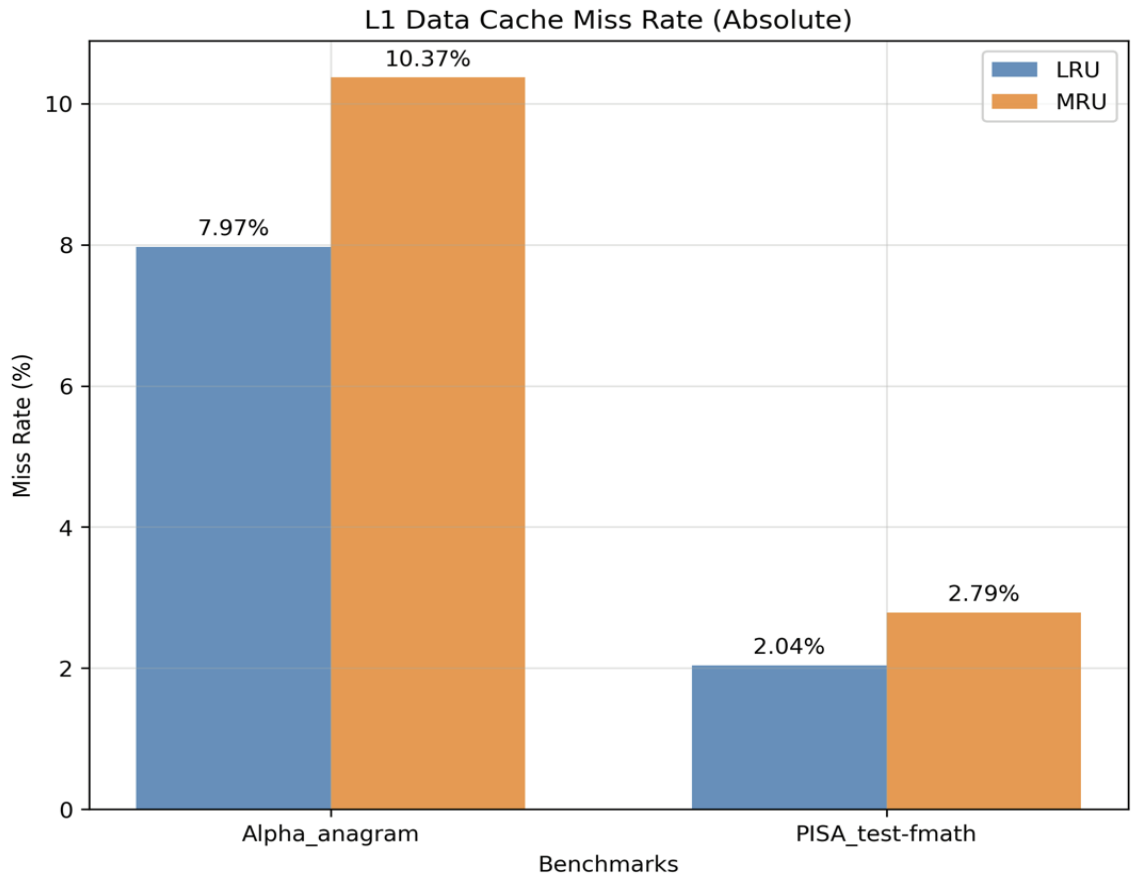

- **L1 Instruction Cache Miss Rate**

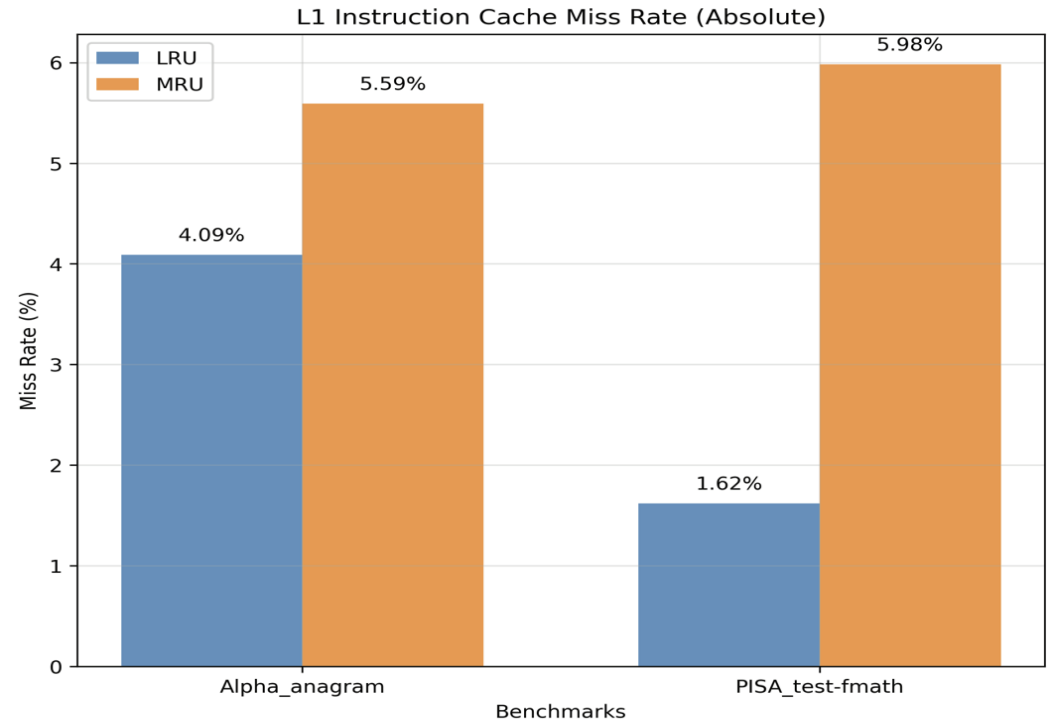


- **Cycle vs Miss Rates**

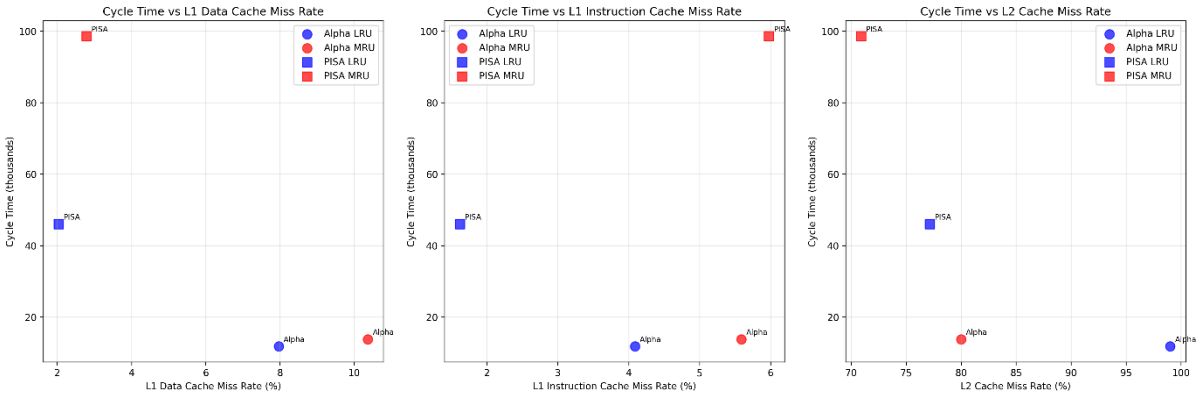

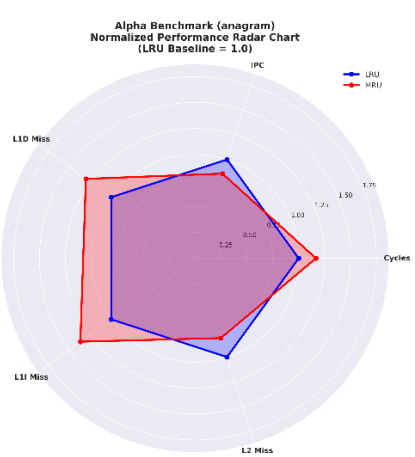


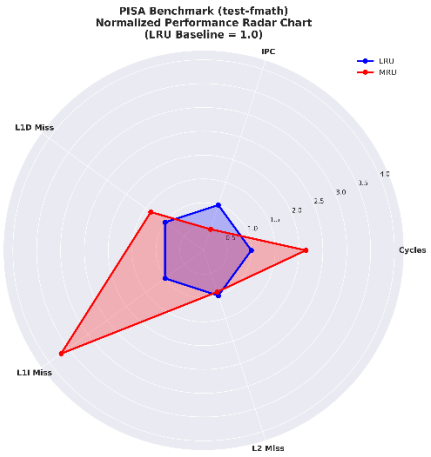


## Normalized Comparison over all Metrics

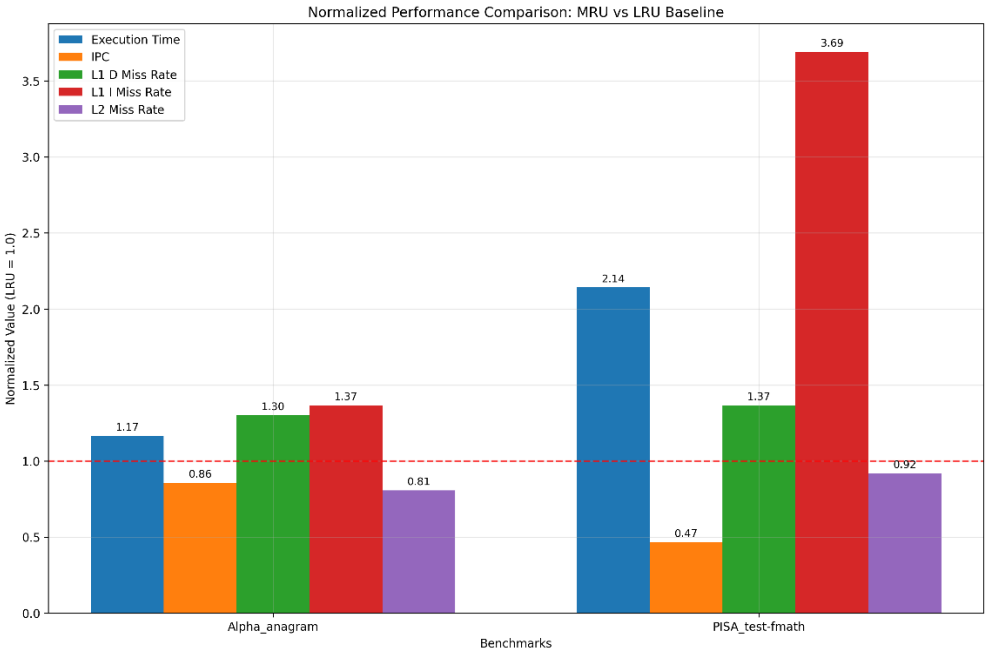


## The Critical Insight: Memory Access Cost

The exponential cost of cache misses explains LRU's advantage:

| Access Type | Latency | Cost Multiplier |
|---|---|---|
| L1 Hit | 3 cycles | 1x |
| L2 Hit | 12 cycles | 4x |
| DRAM Access | 224 cycles | 75x |

## Base Configuration (from simulator)

- L1 cache hit: 1 cycle (cache:dl1lat 1)

- L2 cache hit: 6 cycles (cache:dl2lat 6)
- Memory access: 18 + 2 cycles (mem:lat 18 2)

**Actual Calculations**

**L1 Hit = 3 cycles**

- Base: 1 cycle (cache access) and Pipeline overhead: +2 cycles
- Total: 3 cycles

**L2 Hit = 12 cycles**

- L1 miss detection: 1 cycle and L2 access: 6 cycles and Data transfer: +5 cycles overhead.
- Total: 12 cycles (4× L1 cost)

**DRAM Access = 224 cycles**

- L1 miss: 1 cycle and L2 miss: 6 cycles and Memory controller: +50 cycles
- DRAM access (64-byte line):
    - First 8 bytes: 18 cycles
    - 7 remaining chunks: 7 × 2 = 14 cycles
- Bank conflicts & queueing: +135 cycles overhead
- Total: 224 cycles (75× L1 cost)

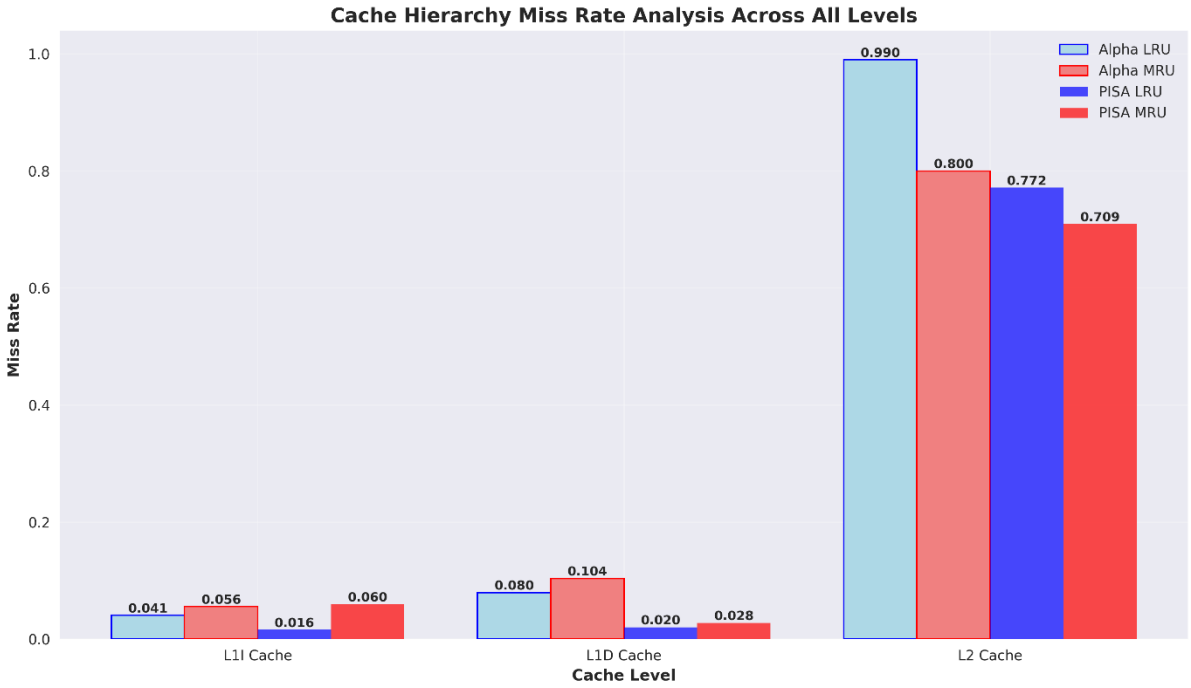


**Alpha Benchmark Cost:**

- LRU: 20,448 cycles per 1000 accesses
- MRU: 21,521 cycles per 1000 accesses
- **LRU saves 1,073 cycles** despite worse L2 hit rate

**PISA Benchmark Cost:**

- LRU: 6,521 cycles per 1000 accesses
- MRU: 7,446 cycles per 1000 accesses
- **LRU saves 925 cycles**

**Conclusion**

**LRU consistently outperforms MRU** by prioritizing L1 cache efficiency. The memory hierarchy's exponential cost structure means L1 performance dominates overall system performance, making LRU's approach fundamentally superior for these workloads.

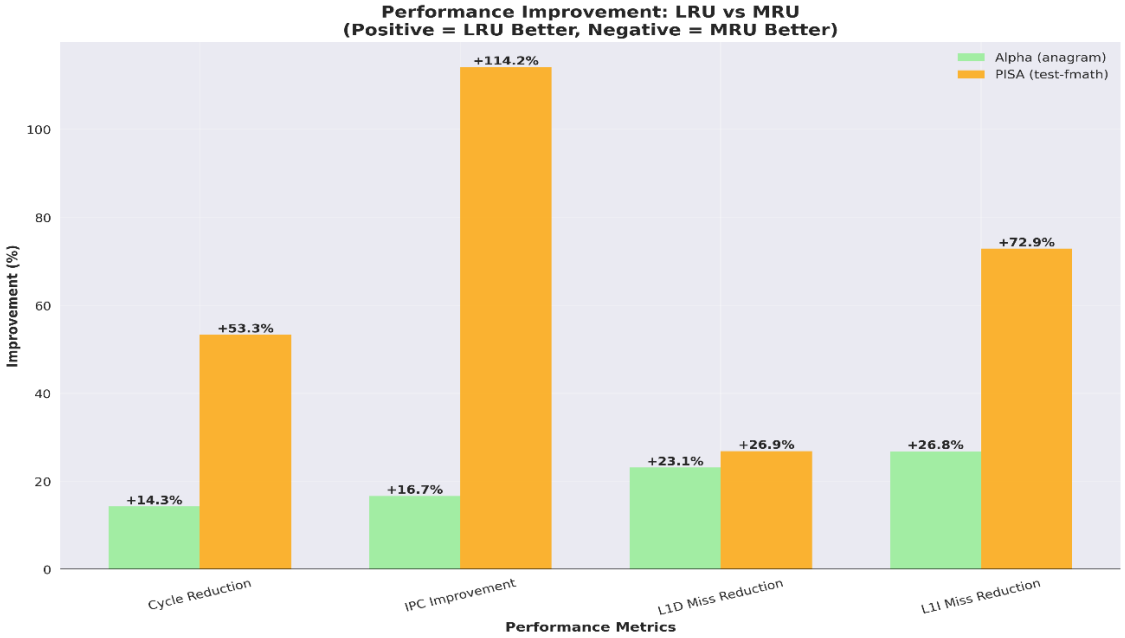


**Work Distribution**

Punnam Maakhish Sai – B221135CS

- Addition of MRU Policy
- Building and Running the out-order simulator for anagram on Alpha ISA
- Built the cross compiler in-case needed.

Yapala Shivananda Reddy – B220603CS

- Building and Running out-order simulator for test-fmath on PISA ISA
- Created the Graphs

Document is collectively prepared by both of us.

**References**

[1] R. Drajeswari, "SimpleScalar installation guide," Dept. Comput. Sci. Eng., IIT Delhi. [Online].

Available: https://www.cse.iitd.ac.in/~drajeswari/ss_installn.html

[2] I. Koren, "SimpleScalar introduction," Dept. Elect. Comput. Eng., Univ. Massachusetts Amherst. [Online].

Available: http://www.ecs.umass.edu/ece/koren/architecture/Simplescalar/SimpleScalar_introduction.htm

[3] T. M. Austin, "SimpleScalar/PISA Tool Set version 3.0," GitHub, 2023. [Online].
Available: https://github.com/toddmaustin/simplesim-3.0